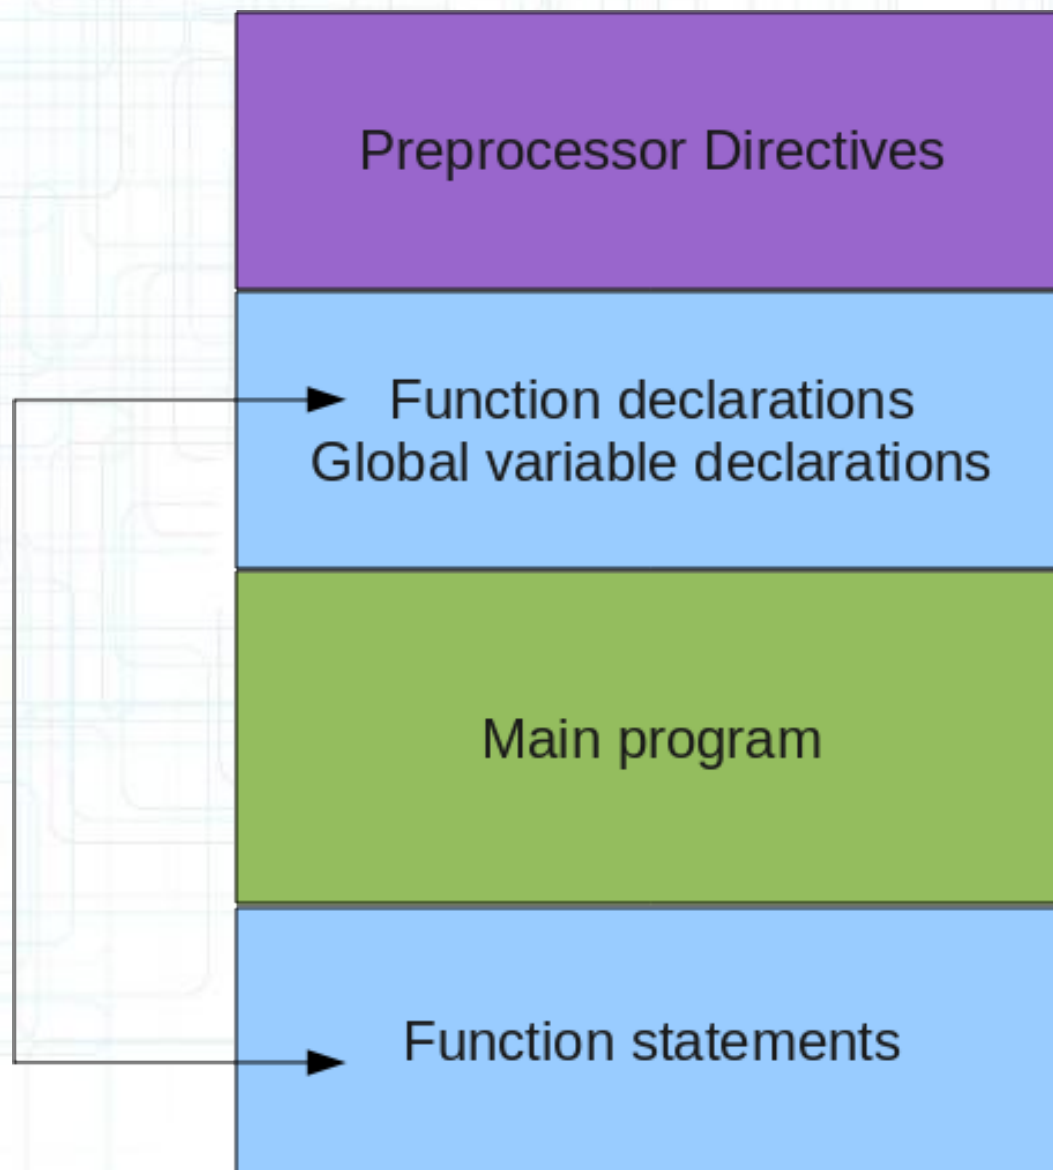


# MPI programming with C

# Structure of a C program revisited



# Modifications on a C source file to implement MPI

- 1 preprocessor directive.
- Function calls on the main program.
- Compile with new compiler command.
- Execute with run-time command.



# How do MPI programs work?

- Basically a copy of the program is executed on every computer (node) in the network (cluster) where it is launched.
- They communicate with each other by sending messages.
- You specify on how many nodes you want it to run.
- Some constraints apply (execution time, number of nodes, no interactivity) when using research center clusters (TACC, Teragrid, etc.).

# Master/Worker paradigm

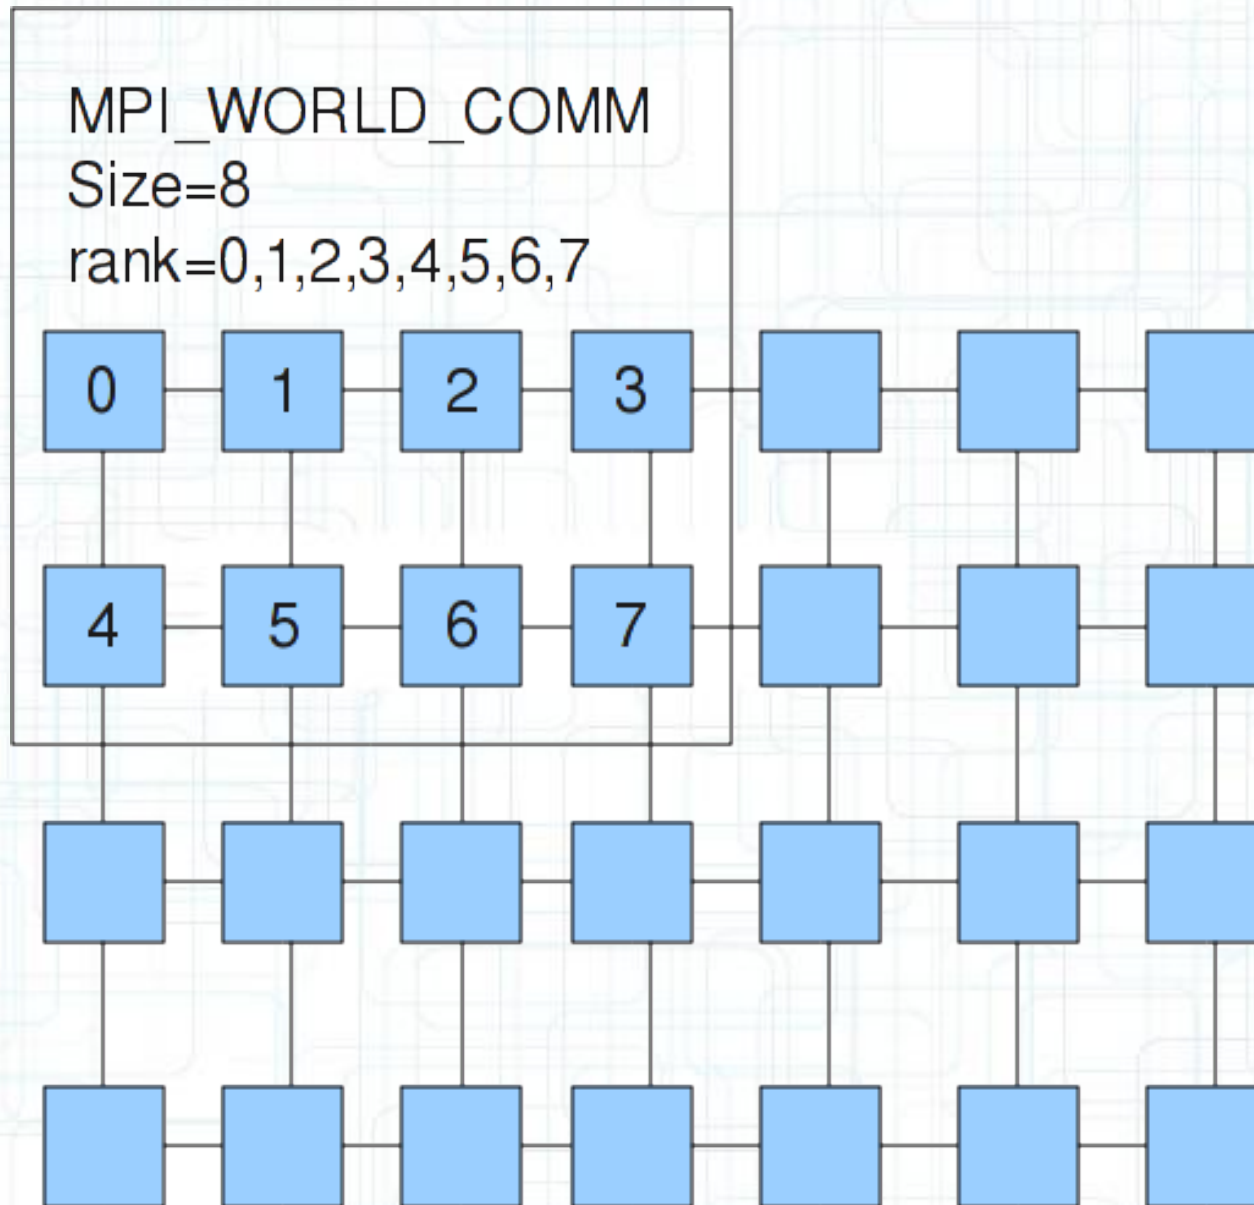
- One node distributes work to others and receives results from them to prepare a final output.
- Serves as an intermediary with the user.
- Reads input and writes output.
- May or may not do work.

# Communicator

- An execution “environment” within the cluster created when a MPI program is executed in it.
- Default name is `MPI_COMM_WORLD`
- Numbers nodes from  $0$  to  $n-1$  for  $n$  nodes that execute the program.
- Node  $0$  is usually the *master* or *root* node.



# Example



# MPI source files

- No need to write a separate program for every node.
- Just make a block of code that corresponds to either the master or worker nodes.
- We differentiate them through their *rank*.



# Sample program

```
int main(int argc, char *argv[])  
{  
    if(rank==0)  
    {  
        Code for master node  
    }  
    else  
    {  
        Code for worker nodes  
    }  
    return 0;  
}
```

# Basic Modifications

- `#include <mpi.h>`
- `MPI_Init(&argc,&argv);`
- `MPI_Comm_rank(MPI_COMM_WORLD,&rank);`
- `MPI_Comm_size(MPI_COMM_WORLD,&size);`
- `MPI_Finalize();`

```
#include <mpi.h>
int main(int argc, char *argv[])
{

    int rank, size;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_WORLD_COMM, &size);
    MPI_Comm_rank(MPI_WORLD_COMM, &rank);
    if(rank==0)
    {

        Code for master node

    }
    else
    {

        Code for worker nodes

    }

    MPI_Finalize();
    return 0;

}
```




# Communication

- Root and worker nodes will communicate by interchanging messages between each other.
- Done with MPI\_Send() and MPI\_Recv() function calls.
- MPI\_Send(&data\_snd,number,MPI\_Type,dest,tag,MPI\_WORLD\_COMM);
- MPI\_Recv(&data\_rcv,number,MPI\_Type,source,tag,MPI\_WORLD\_COMM,&status);

```
#include <mpi.h>
int main(int argc, char *argv[])
{

int rank, size;
MPI_Init(&argc, &argv);
MPI_Comm_size(MPI_WORLD_COMM, &size);
MPI_Comm_rank(MPI_WORLD_COMM, &rank);
if(rank==0)
{
```



```
MPI_Finalize();
return 0;

}
```

# Collective communication

- `MPI_Send()` and `MPI_Recv()` are point-to-point communications, but can be looped to communicate with all nodes.
- There are MPI functions to send or receive messages to and from all nodes in the communicator.
- In order to work every node in the communicator must execute the function call.



# Collective Communication functions

- `MPI_Bcast()`; broadcast a message to all nodes in the communicator.
- `MPI_Reduce()`; get a message from every node in the communicator and do an operation on them.
- `MPI_Scatter()`; distribute an array to every node in the communicator.
- `MPI_Gather()`; fill an array with elements from every node in the communicator.
- `MPI_Barrier()`; set a synchronization barrier.

```

#include <mpi.h>
int main(int argc, char *argv[])
{

int rank, size;
MPI_Init(&argc, &argv);
MPI_Comm_size(MPI_WORLD_COMM, &size);
MPI_Comm_rank(MPI_WORLD_COMM, &rank);
if(rank==0)
{
    Code for master node
}
else
{
    Send and receive
    Code for worker nodes
}

Collective Communication calls

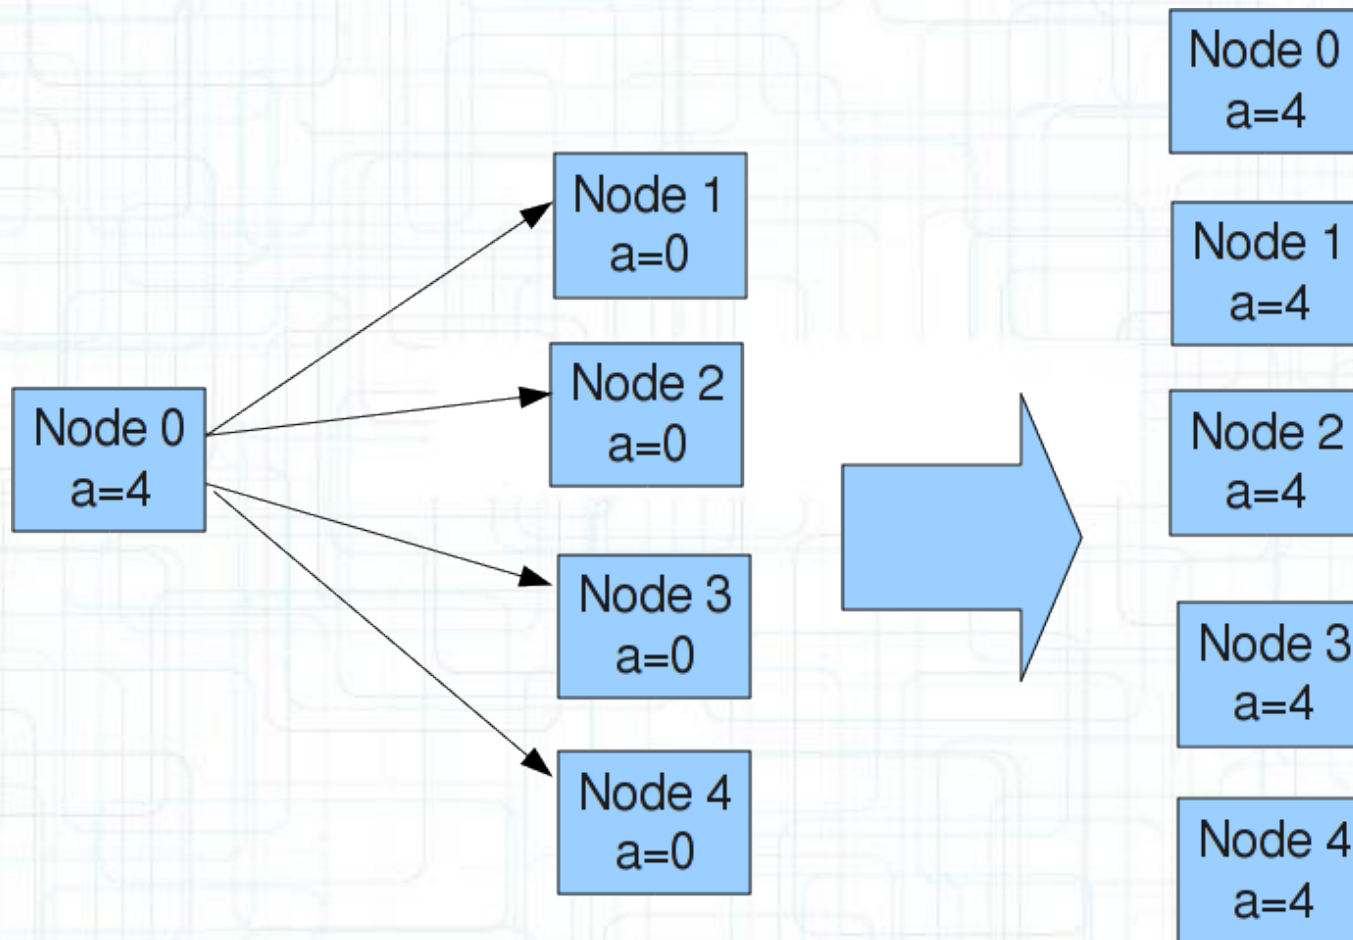
if(rank==0)
{
    Code for master node
}
else
{
    Send and receive
    Code for worker nodes
}

Collective Communication calls

MPI_Finalize();
return 0;
}

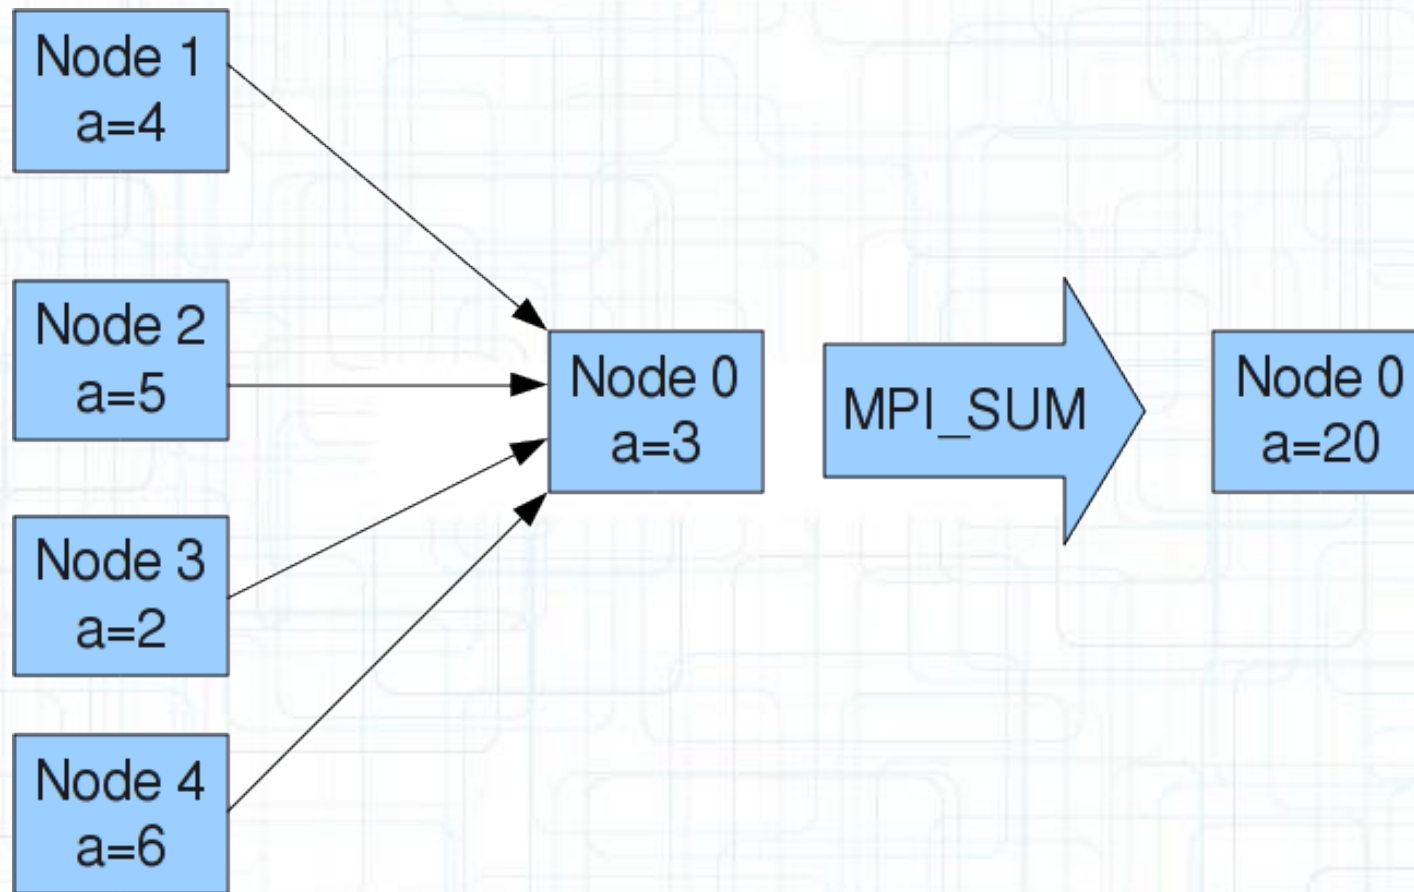
```

# Bcast

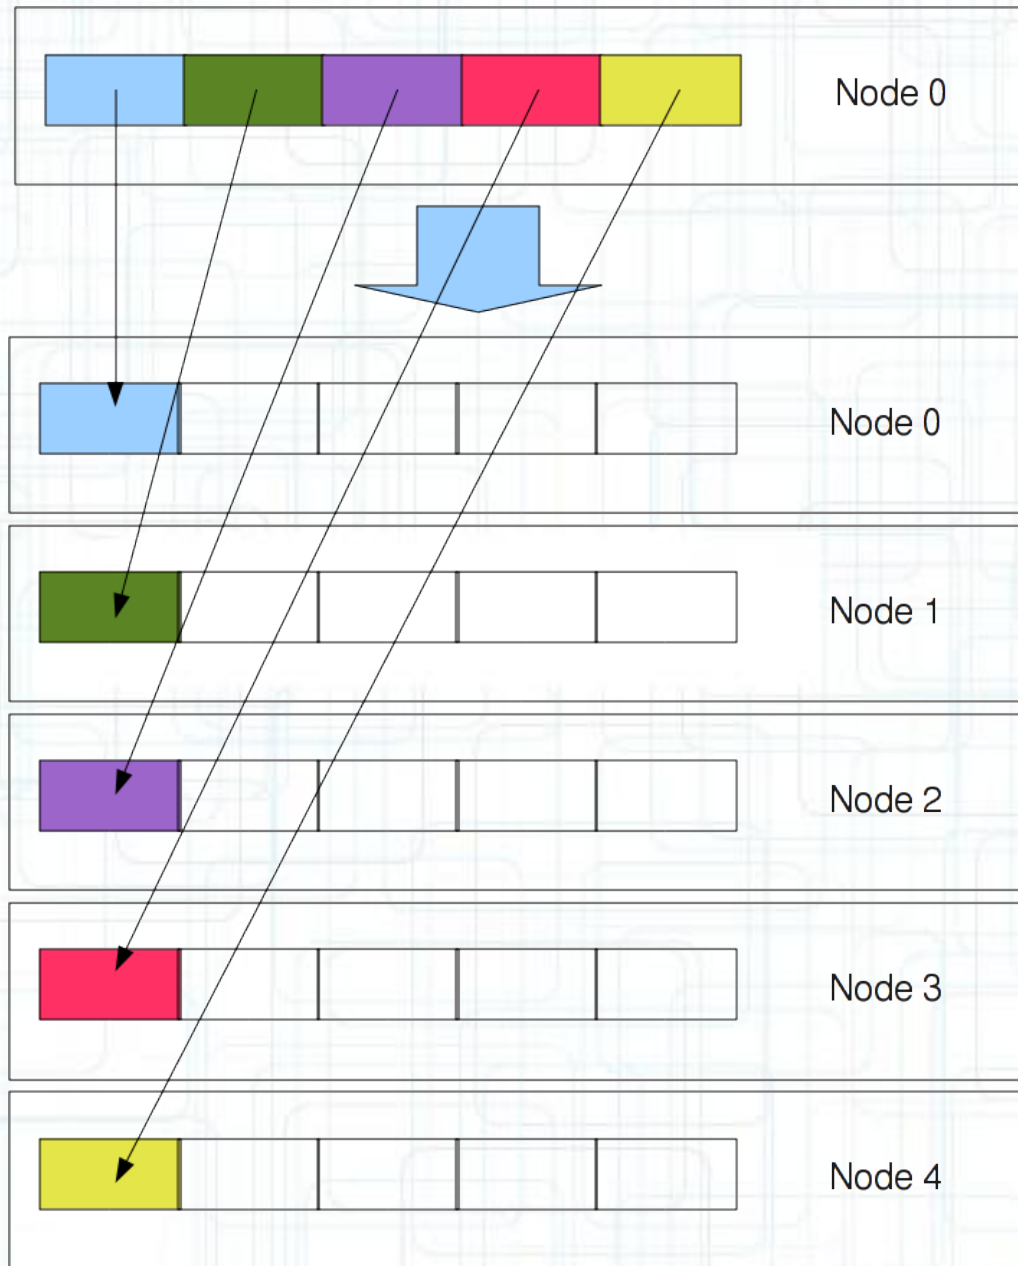




# MPI\_Reduce



# MPI\_Scatter



# Compilation

- for c (`gcc myfile.c -o myfile`)
- for MPI (`mpicc mympifile.c -o mympifile`)
- for c++(`g++ myfile.cpp -o myfile`)
- for MPI(`mpi++ mympifile.cpp -o mympifile`)



# Execution

- `mpirun -n # mympifile`
- `mpiexec -n # mympifile`