

# Collective communications

## Global Reduction Communications

What are these communications ?

Communications that involve computation operation on the data which is transmitted over the processes in the communicator

### **Advantage:**

- collect all the data across the communicator
- reduce multiple data to a single one
- perform desired operation with the data
- store on root or share data over the processes

## MPI Reduce

In C:

```
MPI_Reduce(void *sendbuffer, void *recvbuffer, int count,  
MPI_Datatype data_type, MPI_Op operation, int root, MPI_Comm comm);
```

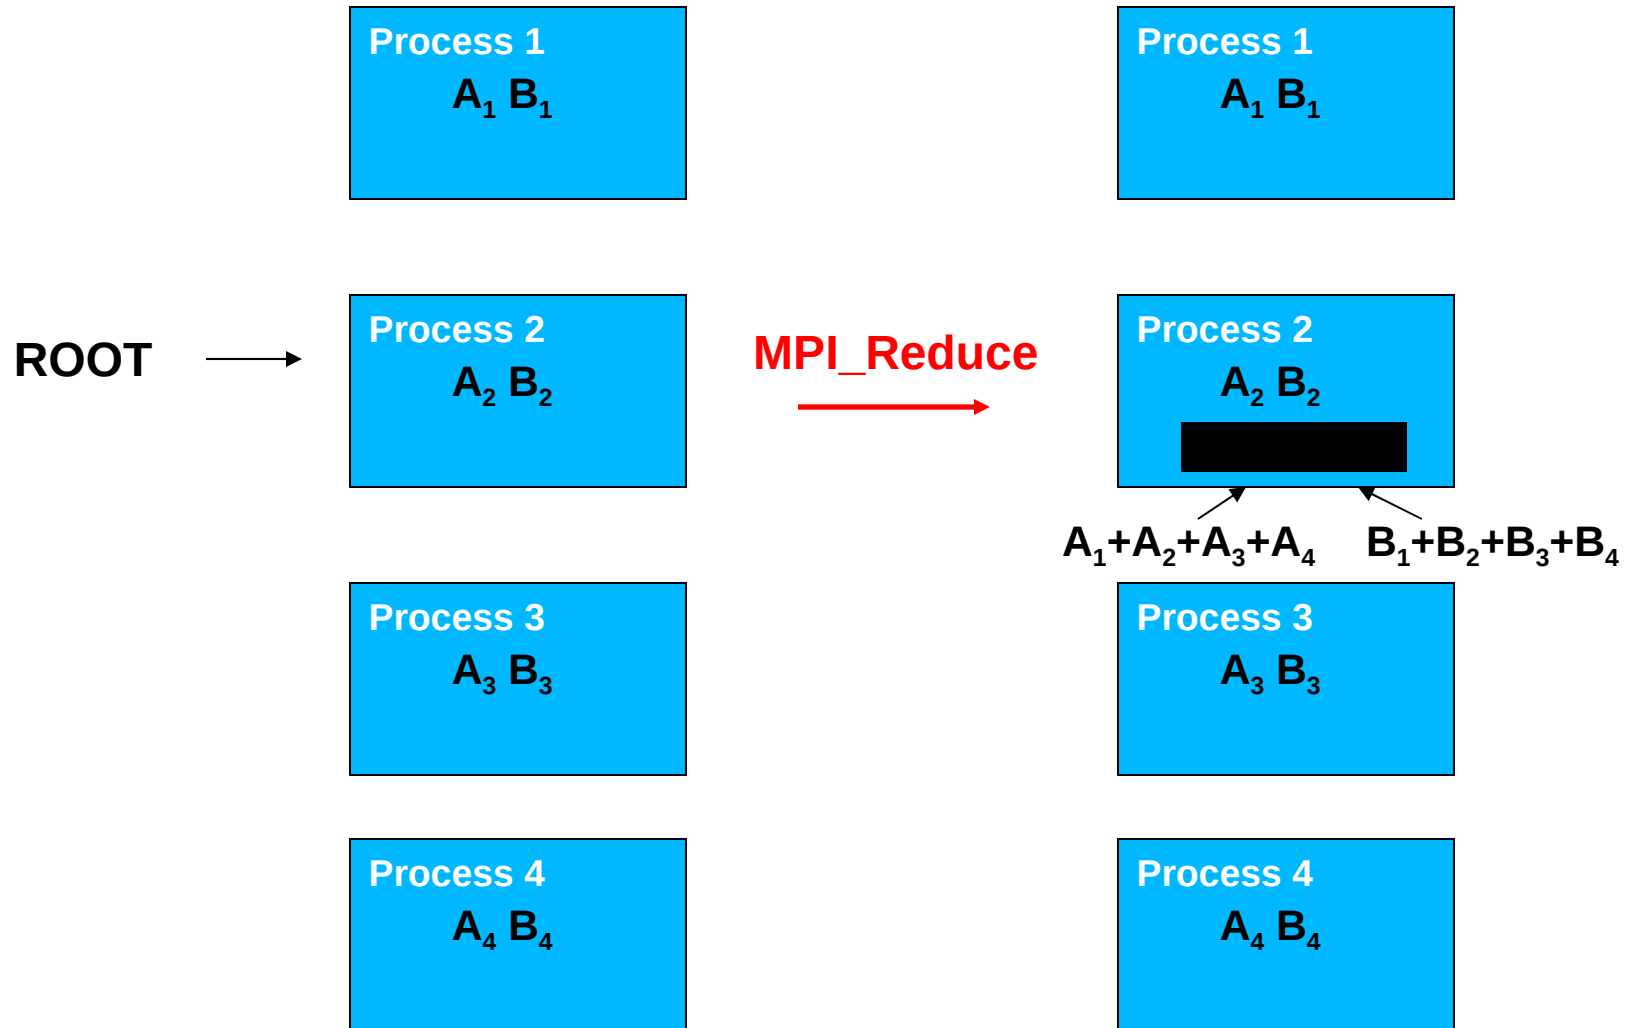
In Fortran:

```
call MPI_Reduce(sendbuffer, recvbuffer, count,
```

# Collective communications

MPI Reduce :





OPERATION = SUMMATION



# Collective communications


MPI Allreduce :

OPERATION = SUMMATION



There is no  
ROOT

**MPI\_Allreduce**



# Collective communications

## Reduction operations

MPI Name	Operation
MPI_SUM	summation
MPI_PROD	product
MPI_MIN	minimum
MPI_MAX	maximum
MPI_LAND	logical AND
MPI_LOR	logical OR

## Other useful reduction subroutines:

MPI\_REDUCE\_SCATTER  
MPI\_SCAN

## Other useful collective subroutines:

MPI\_GATHER, MPI\_ALLGATHER  
MPI\_ALLTOALL

# Collective communications

## Example demonstrating MPI\_Reduce operation: Calculation of PI

Look for “**calculate\_pi.c**”

```
#include "mpi.h"
#include <stdio.h>
#include <math.h>
int main( int argc, char *argv[] )
{
    int n, myid, numprocs, i;
    double PI25DT = 3.141592653589793238462643;
    double mypi, pi, h, sum, x;
    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD,&myid);
    while (1) {
        if (myid == 0) {
            printf("Enter the number of intervals: (0 quits) "); fflush(stdout);
            scanf("%d",&n);
        }
        MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
        if (n == 0)
            break;
```

# Collective communications

## Example demonstrating MPI\_Reduce operation: Calculation of PI

Look for “**calculate\_pi.c**”

```
else {
    h = 1.0 / (double) n;
    sum = 0.0;
    for (i = myid + 1; i <= n; i += numprocs) {
        x = h * ((double)i - 0.5);
        sum += (4.0 / (1.0 + x*x));
    }
    mypi = h * sum;
    MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM, 0,
              MPI_COMM_WORLD);
    if (myid == 0)
        printf("pi is approximately %.16f, Error is %.16f\n",
              pi, fabs(pi - PI25DT));
}
}
MPI_Finalize();
return 0;
}
```

# Collective communications

## Example demonstrating MPI\_Reduce operation: Calculation of PI

Running “**calculate\_pi.c**”

```
mpirun -np 4 ./calculate_pi.exe
```

Output:

```
Enter the number of intervals: (0 quits) 100  
pi is approximately 3.1416009869231249, Error is 0.0000083333333318  
Enter the number of intervals: (0 quits) 500  
pi is approximately 3.1415929869231269, Error is 0.0000003333333338  
Enter the number of intervals: (0 quits) 1000  
pi is approximately 3.1415927369231262, Error is 0.0000000833333331
```